

<p>Main Memory:</p> <ul style="list-style-type: none"> Logical(V) Address: an address generated by software based on the program instructions Physical Address: an address seen by memory unit Base and Limit registers define the legal address space for each process. Multiple-partition allocation: <ul style="list-style-type: none"> Hole – block of available memory Where process arrives, it allocates to memory from a hole large enough OS maintains info (a) allocation partitions (b) free partitions (hole) Dynamic-storage allocation(first-fit, best-fit, worst fit) External Frag: total memory space exists to satisfy a request but is not contiguous Internal Frag: allocated memory may be slightly larger than requested memory Compaction: reduce external frag, shuffle memory, place all free memory in 1 block Page #: index into a page table with base address of each page in physical mem. Page offset: combined with base addr, define physical mem addr is sent to mem unit To run program of n pages, find n free frames and load program <ul style="list-style-type: none"> set up page table to translate logical to physical addresses Frames: fixed-sized blocks Pages: divide logical memory into blocks of same size TBR: points to the page table RLR: indicates size of the page table LB's: special fast-lookup hardware, caches the address translation LB must be consistent with page table Segmentation: memory-management scheme that supports user view of memory Seg Table: map 2D logical addresses to one-dimensional physical addresses <p>Virtual Memory:</p> <p>Virtual Memory: only part of the program needs to be in mem for execution</p> <ul style="list-style-type: none"> separation of user logical/physical memory more efficient process creation address spaces, shared by several processes implemented via (a)demand paging (b)demand segmentation <p>Demand Paging: bring a page into memory only when it is needed</p> <ul style="list-style-type: none"> less I/O needed less memory needed faster response more users <p>Page Replacement: over-allocation of memory due to demand paging</p> <ul style="list-style-type: none"> Sol: replace an existing not currently being used memory with the new one <p>Modify(dirty) bit: reduce overhead of page trans, only modified page write to disk</p> <p>FIFO page algorithm: replace the oldest page</p> <p>Optimal Page Algorithm: replace page that will not be used for long periods of time</p> <p>Least Recently Used: replace page that has not been used for the longest period</p> <p>Thrashing: a process is busy swapping pages in and out</p> <ul style="list-style-type: none"> occurs when size of locality > total memory size <p>Slab Allocator:</p> <ul style="list-style-type: none"> slab: one or more continuous pages cache: consists of 1 or more slabs, single cache for each unique kernel structure when a new object is needed, the slab allocator tries to find a “free” object; if found, it is marked as used. Benefits: no fragmentation, fast memory when objects are frequently allocated & deallocated <p>File System:</p> <p>File: a named collection of related information that is recorded in secondary storage.</p> <p>File Ops: create, read, write, reposition, delete, truncate</p> <p>Open(Fi): search for Fi and move content to memory</p> <p>Close(Fi): move content of entry Fi to directory</p> <p>Disk Structure:</p> <ul style="list-style-type: none"> partitions raw (unformatted) or formatted with a file system volume: any entity containing file system general purpose file systems special purpose file systems <p>Single Level Directory: single directory for all users, name problem.</p> <p>Two-Level Directory: separate directory for each user</p> <ul style="list-style-type: none"> path name defined by a user name and a file name <p>Tree-Structured Directories:</p> <ul style="list-style-type: none"> absolute path name: begins at the root, follows a path down to specified file Relative Path Name: defines a path from the current directory <p>File-system Structure: file system resides on secondary storage</p> <ul style="list-style-type: none"> logical file system file-organization module basic file system I/O control <p>on-disk File System Structures:</p> <ul style="list-style-type: none"> mount table: info about each mounted volume directory-structure cache: directory info of recently accessed directories system-wide open-file table: copy of FCB of each open file Per-Process open-file Table: a pointer to an entry in sys-wide open-file table Buffers: file-system blocks when read from or written to disk <p>A file system must be mounted before it can be accessed</p>	<p>Virtual File System: abstraction of generic file system operations from specific file system implementations.</p> <p>Four main object types:</p> <ul style="list-style-type: none"> Inode: object represents an individual file File: an open file Superblock: entire file system Dentry: an individual directory entry <p>Directory Implementation</p> <ul style="list-style-type: none"> Linear List: of file names with pointer to the data blocks Hash Table: linear list with hash data structure <p>Allocation Methods:</p> <ul style="list-style-type: none"> Contiguous (sharing a common border) Allocation: each file occupies a set of contiguous blocks on the disk Simple: only starting location (block#) and length(num blocks) required Random access is easy Waste of space Files cannot grow <p>Linked Allocation: each file is a linked list of disk blocks which may be scattered anywhere on the disk</p> <ul style="list-style-type: none"> the directory contains a pointer to the first and last blocks of a file simple: need only starting address no wasted space no random access space overhead for storing the pointers reliability problem caused by lost/bad pointers <p>File Allocation Table: one entry for each block and is indexed by block number. Each block's table entry contains a block number of the next block in the file</p> <ul style="list-style-type: none"> improved random access: disk head can find the location of any block by reading FAT. FAT can be cached to reduce disk head movement. Greater space overhead than linked allocation <p>Indexed Allocation: handle large files, each file has its own index block, which is an array of disk=block addresses.</p> <ul style="list-style-type: none"> support random access greater space overhead than linked allocation <p>Combined Scheme in UNIX</p> <ul style="list-style-type: none"> inode keeps 15 pointers first 12: point to direct blocks next 3 to indirect blocks (a)single indirect block(b)double(c)triple <p>Free-Space Management: keep track of free space for reuse</p> <p>Linked Free Space List on Disk:</p> <ul style="list-style-type: none"> cannot get contiguous space easily no waste of space <p>Mass Storage Structure:</p> <ul style="list-style-type: none"> magnetic disks (secondary storage) drive attached to computer via I/O bus <p>Disk Drives: addressed as large 1-dimensional arrays of logical blocks</p> <p>Low-level(Physical) Formatting: divide disk into sectors that the disk controller can read and write.</p> <p>Partition: disk into one or more groups of cylinders</p> <p>Logical Formatting: “creating a file system”, store initial file system data to disk</p> <p>Access Time:</p> <ul style="list-style-type: none"> Seek time: time for the disk arm to move the heads to the cylinder containing the desired sector. (Seek time = seek distance) Rotational Latency: additional time waiting for the disk to rotate the desired sector to the disk head. <p>Disk Scheduling</p> <ul style="list-style-type: none"> FCFS: fair but not fastest, total head movement 640 cylinders SSTF: select request with minimum seek time, may cause starvation, not optimal, 236 cylinders SCAN (elevator Algorithm): disk arm starts at one end and moves toward the other, servicing. Has an uneven wait time C-SCAN: when reaches the other end, it immediately returns to the beginning of the disk, without servicing any requests on the return trip <p>Swap-space Management: virtual memory uses disk space as an extension of main memory.</p> <ul style="list-style-type: none"> kernel uses swap maps to track swap-space use <p>RAID Structure:</p> <ul style="list-style-type: none"> redundant arrays of independent (inexpensive) disks improve reliability through redundancy, duplication disks improve performance through parallelism, can read multiple disks <p>RAID Levels:</p> <ul style="list-style-type: none"> 0: non-redundant striping 1: mirrored disks 2: memory-style error correction codes 3: bit-interleaved parity 4: block-interleaved parity 5: block-interleaved distributed parity <p>Solid-State Drive: storage device using memory to store persistent data</p> <ul style="list-style-type: none"> higher reliability
---	---

- consistent read performance
- more power efficient

Disadvantages:

- more expensive
- less capacity
- less write longevity
- slower write speeds

I/O Hardware:

- Port: device communication with computer
- Bus: set of wires, rigidly-defined messages that be sent on the wires
- Controller: electronic that can operate a port, a bus (host adapter), or device

CPU talks to device controllers through registers.

Polling: determines state of device

- command-ready
- busy
- error

Interrupts:

- interrupt-request line: triggered by I/O device
- Interrupt Handler: receives interrupts
- Maskable: ignore or delay some interrupts
- Interrupt Vector: dispatch interrupt to correct handler
- Also used for exceptions (divide by zero)

Direct Memory Access:

- used to avoid programmed I/O for large data movement
- requires DMA controller

Perform DMA Transfer:

1. device driver is told to transfer disk data to buffer at address X
2. device driver tells disk controller to transfer C bytes from disk to buffer at X
3. disk controller initiates DMA transfer
4. disk controller sends each byte to DMA controller
5. DMA controller transfers bytes to buffer X, increasing memory address, dec C
5. When C = 0, DMA interrupts CPU to signal transfer completion

Application I/O Interface: IO calls encapsulate device behaviors in generic classes

Programmable Interval Timer: used for timings, periodic interrupts, always a copy

Kernel I/O Subsystem:

- caching: fast memory holding copy of data, always a copy
- spooling: hold for a device, if device can serve only 1 request at a time
- device reservation: exclusive access to a device, system calls for alloc & deal

Protection: ensure that OS object is accessed correctly and only by those processes that are allowed to do so.

Domain Structure:

- rights-set: subset of all valid operations that can be performed on the object
- Domain is a set of access-rights associated with an subject (entity), can be switched via password, file system, command

Access Matrix:

- View protection as a matrix (access matrix)
- rows represent domains
- columns represent objects
- access(i, j): set of operations that a process executing in domain i can invoke on object j
- Access control can be discretionary or mandatory
- DAC: user who creates object can define access column for that object
- MAC: sys admin determines the access matrix, user cannot modify it
- Use of Access Matrix: can be expanded to dynamic protection
- owner: return the owner of oi
- copy: copy op from oi to oj
- control: di can modify dj access rights
- transfer: switch from domain di to dj
- mechanism: ensures matrix can be manipulated by authorized entities
- policy: user define the matrix: who can access what object in what mode
- Implementation of Access Matrix: global table, per-object access list, per-domain capability list, lock-key.
- For access matrix each column defines an access-control list for the object: domain, operation
- Each row defines a capability list (like a key): object, operation

Revocation of Access Rights:

Immediate vs. delayed, selective vs. general, partial vs. total, temp vs. permanent

Security:

Threat: potential security violation

Attack: attempt to breach security

Security Violation Categories:

- Confidentiality: unauthorized reading of data
- Integrity: unauthorized modification of data
- Availability: unauthorized destruction of data

Security Violation Methods:

- Theft of service: unauthorized use of resources
- Denial of service (DOS): prevention of legitimate use
- Masquerading: pretending to be an authorized user

Man-in-the-middle attack: intruder sits in data flow, masquerading as sender to receiver and vice versa

Session hijacking: intercept an already-established session to bypass authentication

Security Measure Levels:

- physical: data centers, servers, connected terminals
- host: operating system
- network: intercepted communications, interruption, DOS
- human: social engineering, phishing, dumpster diving

Virus: code fragment embedded in legitimate program

Port Scanning:

Port scan: automated attempt to connect to ports on one or more hosts

goal: detection of answering service protocol

Denial of Service: overload target computer preventing it doing any useful work

Cryptography: symm encryption/asymm encryption (public/private key)

Firewall and Network Security:

A firewall is placed between trusted and untrusted hosts

A firewall limits network access between these two security domains

Main Memory:

Main memory and registers are only storage CPU can directly access

linker (or loader) binds relocatable addresses to absolute addresses

Address Binding: compile, load, or execution time

OS determines logical to physical address mapping using MMU.

MMU is hardware that maps logical address to physical address at run time

MMU converts virtual address to physical address

segmentation and paging are two typical types of MMU

Paging has no external fragmentation, but internal fragmentation

Associative memory: memory that supports parallel search

hit ratio – percentage of page translation that is found in the TLB

Page Sharing

- one copy of code shared by all processes of the same program
- shared memory can be used for inter-process communication

To reduce memory consumption of page tables:

- hierarchical page table, hashed page table, inverted page table

Hierarchical Page Tables: Break logical address space to multi-level of page tables

Two-Level Paging:

A logical address is divided into:

- a page directory number (first level page table)
- a page table number (2nd level page table)
- a page offset

Hashed Page Tables: each element contains: page#, frame#, pointer next element

Inverted Page Table -entry has process id and the page# (virtual address)

Segmentation: supports user view of a program a program is a collection of segs

- main program
- function
- local variables, global variables
- stack
- each segment can be mapped to physical blocks

Segment table maps segments to physical memory

each segment table entry has:

- base: the starting physical address where the segments reside in memory
- limit: the maximum offset of the segment
- memory protection bites: present/read/write/execution

Swapping extends physical memory with backing disks

a process can be swapped temporarily out of memory to a backing store

backing store is usually a (fast) disk

the process will be brought back into memory for continued execution

Threads:

A thread is an independent stream of instructions that can be scheduled to run as such by the kernel

Benefits: responsiveness, resource sharing, economy, scalability

Signals are used in UNIX systems to notify a process that a particular event has occurred.

synchronous signals are delivered to the same thread causing the signal

Dispatcher module gives control of the CPU to the process selected by the short-term scheduler

System call is a programming interface to access the OS services

Process is a program in execution, its execution must progress in sequential fashion. a program is static and passive, process is dynamic and active

Long-Term Scheduler (Job Scheduler)- selects which processes should be brought into the ready queue

Short-Term (CPU Scheduler)- selects process to be executed next & allocates CPU

Mid-term Scheduler- swap in/out partially executed process to relieve memory pressure

